

Models Take Notes at Prefill: KV Cache Can Be Editable and Composable

Bojie Li
Pine AI

Code: <https://github.com/bojieli/programmable-kv> Interactive companion:
<https://01.me/research/programmable-kv/>

Abstract

Prefix caching reuses a transformer’s prefill only across an exactly shared prefix, so a single changed field—a clock tick, a session id, an account-status flip—invalidates the entire downstream cache. We start from a puzzle: the cache *before* a field is provably reusable (key/value deviation 0.0), yet surgically overwriting only the field’s own key/value vectors and reusing the rest leaves the model acting on the *old* value.

The reason—established causally by four probes, replicated across four model families—is that at prefill the model has already computed the *field-conditioned conclusion* and written it onto downstream aggregator/delimiter tokens: the field’s own key/value drives under 1% of the decision, while the downstream *notes* drive essentially all of it. The KV cache is best read as a notebook of memoized conclusions, and two capabilities follow. **It is editable.** Because the conclusion is already written downstream, the cheap, robust fix is a salient *erratum* that amends the notes (`field+erratum` matches the strong hoist-to-end baseline without rewriting the prompt); one can instead recompute the affected downstream notes—reliably for the whole affected suffix, but only unreliably for a cheap top- K (`field+selective@K`, with a model-dependent minimum). An even cheaper in-place refresh of the field alone is gated by *chain-of-thought*, not model size: with CoT the chain re-reads the field and the edit alone recovers the decision (1.00 at 8B, $\sim 1\%$ compute), without CoT it is ignored (0.00). **It is composable.** The notes are localized and position-portable, so a precompiled skill can be RoPE-repositioned and spliced into any context—behaviorally indistinguishable from full recompute (logit cosine 0.90–0.999 across twelve models) at $O(L)$ rather than $O(L^2)$ time-to-first-token ($13.9\times$ at 32k). A keystone experiment—editing a field *inside* a transplanted skill—reproduces the editing mechanism verbatim, showing the two operations act on one substrate; a unified edit+compose agent (thirteen models, up to 300 decisions each) stays decision-identical to full recompute (agreement 0.81–1.0) at up to $14.9\times$ lower latency.

The substrate is any per-token attention KV cache. We validate it across scale, quantization, Mixture-of-Experts, and multimodal image caches; extend it to Multi-head Latent Attention, sliding-window, and interleaved-M-RoPE models through small adapters; and map where it holds up to the 2026 sparse-attention frontier. The caching machinery itself is prior art; our contributions are the mechanism, the unification, a decision-governance evaluation, and the attention-variant adapters.

1 Introduction

Modern LLM agents re-read long, mostly-static instructions on every turn—a system policy, a tool specification, retrieved documents. Key/value (KV) caching makes this affordable by reusing the prefill across turns, but only across an *exact* shared prefix. The moment one token changes inside the reused region—a timestamp, a user id, an order’s status—the keys and values of *every* later token are invalidated, because each attended to the token that changed. The de-facto workaround—*hoisting* all mutable content to the end so the static prefix stays cache-aligned—pushes an inference-layer constraint into the application layer: fields referenced in several places, nested sub-agent prompts, and dynamically assembled contexts cannot all be cleanly hoisted, and the application must enumerate every mutable field in advance.

This paper starts from a concrete puzzle. The region of the cache *before* a field is, by construction, independent of the field’s value—we measure a key/value deviation of exactly 0.0 when the field changes. One might therefore hope to *surgically* refresh only the field’s own keys and values, leave the rest of the cache stale, and pay almost nothing. We find this fails completely: the model’s decision reverts to the *old* field value, as if the edit never happened (Figure 1b).

The discovery. The reason, which we establish causally, is that transformers do not defer their reasoning to decode time. At *prefill* the model already computes the *field-conditioned conclusion* and writes it onto downstream tokens—disproportionately onto aggregator/delimiter tokens that later positions attend through. The decision then reads these *notes*, not the field itself: across models the field’s own KV causally drives less than 1% of the decision, while the downstream notes drive essentially all of it. The KV cache is best understood not as a frozen byproduct of prefill but as a *notebook of memoized conclusions* (Figure 1a).

Two capabilities from one mechanism. This reframing is generative. (1) If the conclusion is already written downstream, then *editing* a field means amending the notes, not recomputing them. The cheap and robust fix is a one-line salient *erratum* that overrides the stale notes; one can instead recompute the affected downstream notes (reliably for the full affected suffix, cheaply-but-unreliably for a top- K , *field+selective@K*). An even cheaper in-place refresh of the field alone is gated by *chain-of-thought*: with a reasoning chain the model re-reads the field and the edit suffices, without it the edit is ignored. (2) If the notes are localized and position-portable, then a reusable skill can be *composed* into a new context by repositioning and splicing its cached notes—no recompute. A *keystone* experiment, editing a field *inside* a transplanted skill, shows the two operations act on a single substrate. Figure 2 previews both capabilities.

Contributions.

- **A mechanism** (Section 3)—*attention-mediated memoized inference*—established by four causal probes (locality patching, suffix concentration, linear probing, circuit knockout) and four further controls (content/conclusion dissociation, layer-timing, specificity, note-injection; Appendix C), replicated across *four model families* (Qwen3, Llama-3.1, Gemma-2, Mistral), resolved to a *component-level circuit*—named read/write heads, a causal conclusion direction, an SAE feature, attention-vs-MLP, and causal scrubbing (Appendix D)—and connected to the delimiter-token aggregation seen in interpretability.
- **An editing capability** (Section 4): naive KV editing fails; the *erratum / field+erratum* fix matches the hoist-to-end oracle without prompt surgery; an analysis of when the $\sim 1\%$ -compute in-place edit suffices (it requires reasoning and is strongly model-dependent); and a head-to-head with *weight* editing (ROME, LoRA) showing it is the wrong tool for mutable per-request state (global contamination, collateral, 30–50 $\times$ slower).




Figure 1: **Models take notes at prefill.** (a) At prefill the model memoizes the field-conditioned conclusion onto downstream aggregator tokens (orange); at decode the decision reads those notes (blue). (b) Consequently, surgically editing the field’s own KV is ignored (without a reasoning chain), but the decision is recovered cheaply: recompute the *affected* downstream suffix, or—cheaper and robust—append a salient *erratum* ($P(\text{new decision})$, no-CoT model; with chain-of-thought the field-only refresh alone also works, Figure 2c). (c) Because the notes are position-portable, a precompiled skill can be pasted into a new context in $O(L)$ rather than $O(L^2)$ time.

- **A composing capability** (Section 5): position-portable transplant of precompiled skills, $O(L)$ vs. $O(L^2)$ TTFT ($13.9\times$ at 32k), with a seam-repair knob—honestly credited to prior caching work, with our addition being the mechanism that explains it and a correctness lens.
- **The unification** (Section 6): the keystone and a unified edit+compose agent over thirteen models.
- **An application: user memory** (Section 7): the large, mutable user-memory document is both composed (precompiled, repositioned, spliced) and edited (in-place / erratum) as one substrate—decision-faithful to full recompute at $2.3\text{--}4.3\times$ lower time-to-first-token, validated to 70B and on real long-conversation memory (LoCoMo, transplant $\equiv$ full recompute in QA accuracy)—with a pre-registered, statistically-controlled evaluation.
- **Reach and scope** (Section 8): validation across scale, MoE, quantization, and multimodal image caches; small adapters for MLA and interleaved M-RoPE; a fix for sliding-window; and an analysis of where the substrate holds up to the 2026 sparse/compressed-attention frontier.
- **Systems payoff** (Section 9): a real agentic environment and a comprehensive online vLLM serving benchmark (V1 engine, continuous batching, Poisson load)—98.5% vs. 1% prefix-cache hit-rate, $53\text{--}398\times$ lower p_{90} TTFT, and throughput gains that grow with load to $14.5\times$.

Roadmap. Section 3 establishes the mechanism; Sections 4 to 6 derive editing, composing, and their unification; Section 7 applies both to user memory; and Sections 8 and 9 map the substrate’s reach and the end-to-end systems payoff.




Figure 2: **Editing and composing, previewed** (Qwen3-8B unless noted; detail in Sections 3, 4 and 6). (a) *KV editing landscape*: naive edits (stale, field-only without CoT, CacheBlend) fail, while the append-only erratum/field+erratum reach full-reprefill correctness cheaply and robustly (hoisting also works but needs prompt surgery). (b) *Recompute the affected notes*: recovery climbs as more of the post-field *affected suffix* is recomputed—reliable in the limit, but the cheap top- K version (field+selective@ K) is unreliable (Figure 5). (c) *Chain-of-thought, not model size, gates the cheapest edit*: with CoT the near-free field-only refresh alone recovers the decision (1.00); without CoT it is ignored (0.00); the erratum fixes both. (d) *Composing + editing is lossless*: the unified agent stays decision-identical to full recompute across models.

2 Related work

Where computation is stored, and how to edit it. A line of interpretability work localizes stored *knowledge* in transformer weights and edits it there: ROME and MEMIT [16, 17] locate and rewrite factual associations in MLP weights, while circuit analyses such as the indirect-object-identification study [24] trace how specific computations are carried by attention heads. Weight editing targets *durable, global* facts; we compare against a faithful ROME and a LoRA fine-tune empirically (Table 1) and find them ill-suited to *mutable per-request* state—a global edit contaminates concurrent requests and damages unrelated decisions—which is exactly the niche the editable KV cache fills. Closest in spirit, Lindsey et al. [11] find models commit *plans* to specific tokens (e.g. a planned rhyme stored on a line-break token) during the forward pass. We study a complementary object: not weights and not decode-time computation, but the *KV cache*—the activations an inference system already stores and an editor can directly manipulate—and show it holds *memoized conclusions* concentrated on aggregator/delimiter tokens. To our knowledge this is the first causal account of why in-place KV field editing fails and what to do instead.

KV reuse and composable caching. Reusing precomputed KV beyond an exact prefix is an active systems topic, and our *composing* capability builds directly on it; we claim none of the caching

machinery as novel. Prompt Cache [5] precomputes reusable prompt modules with position placeholders and splices them. CacheBlend [27] reuses non-prefix chunk KV and selectively recomputes $\sim 15\%$ of tokens to restore cross-attention. EPIC [6] introduces position-independent caching with ATTNLINK, recomputing only a few chunk-boundary tokens (exploiting attention sinks) for near-linear recompute. CacheSlide [12] reuses KV in a position-aware way via relative-position-dependent caching, and MPIC [31] extends position-independent caching to the multimodal setting by recomputing image-boundary tokens; KVLink [26] is a further reuse system. Mapped onto our terms, our RoPE-repositioning is CacheSlide’s relative-position reuse, our seam-repair is the boundary recompute of CacheBlend/EPIC/MPIC, and our image-KV transplant is MPIC’s idea (we *re-rotate* M-RoPE rather than recompute). Our contributions over this line are orthogonal: (i) the *mechanism* that explains *why* boundary recompute is what is needed; (ii) a *decision-governance* evaluation—does a transplanted skill still *govern the tool decision*, rather than only preserve perplexity or throughput; (iii) the *editing* axis and the keystone unification; and (iv) adapters that extend the operations to new attention representations (MLA, interleaved M-RoPE, sliding-window). We note honestly that this editable/composable-cache direction grew directly out of the position-independent and position-aware caching of EPIC and CacheSlide [6, 12], and out of discussions with Junhao Hu (first author of EPIC; see Acknowledgements); our additions are the mechanism, the editing axis, and the decision-governance lens, not the reuse machinery itself.

Prefix caching, KV compression, and reuse systems. Production prefix caching (vLLM Automatic Prefix Caching, SGLang RadixAttention) reuses exact prefixes. A large literature instead *compresses* or *evicts* the cache: StreamingLLM keeps attention sinks [25]; H2O [30], Scissorhands [14], and SnapKV [10] evict low-importance tokens; Quest [21] keeps all tokens but attends sparsely. Serving systems stream or share cached KV across requests—CacheGen [13] for fast loading and RAGCache [8] for retrieval reuse. These change *which* tokens are present, and compose with our edit/transplant operations only over retained tokens; we treat them, the latent/decoupled-RoPE representation of Multi-head Latent Attention [3, 4], and the 2026 sparse/compressed-attention designs as scope boundaries in Section 8. Our repositioning relies on rotary position embeddings [20] and their extensions [18].

Activation-level interventions. Beyond weight editing, a body of work intervenes on *activations*: steering and inference-time interventions [9], task and function vectors [7, 22], and the causal-mediation/activation-patching methodology we adapt [23]. These edit residual-stream directions to change behavior; we instead read and write the *KV cache* itself—the per-token activations a serving system already persists—which is what makes the intervention both interpretable and deployable.

Agents and tool use. Our evaluation targets tool-using agents in the style of ReAct [28] and Toolformer [19], and we measure end-to-end task success on the τ^2 -bench tool-agent-user benchmark [2], the dual-control successor of τ -bench [29], where state changes mid-trajectory make cache staleness a first-class correctness problem.

3 The discovery: memoized inference in the KV cache

Setup. We study a minimal but agent-realistic decision: a context contains a policy rule and a *mutable field* whose value determines the correct action (e.g. “cancel an order only if its status is pending”; with `status=shipped` the correct action flips from *cancel* to *deny*). After prefilling the full context we compare four cache states at the decision token: *stale* (old field, old downstream), *field-only* (refresh the field’s KV, leave the downstream stale), *full-downstream* (recompute everything after the field), and *oracle* (a clean prefill of the new value). We report *decision recovery*: the fraction of the oracle’s flip that a cache state reproduces (0 = behaves like stale, 1 = behaves like

oracle). Four independent probes converge on one account (Figure 3); Appendix B gives a complete worked example—the verbatim prompt, the erratum, and recorded model responses.

(1) Locality: the field’s own KV barely matters. Refreshing only the field’s KV recovers essentially none of the decision—field-only recovery is -0.028 on Llama-3.1-8B and near zero across models—whereas recomputing the downstream recovers it fully (1.0) (Figure 3a). The field is read *indirectly*: its causal contribution to the decision is under 1%.

(2) Suffix concentration: the effect lives downstream and late. Sweeping how many downstream tokens we patch (ranked by causal effect) shows recovery accrues slowly and saturates only as we include many tokens after the field, with the causal mass concentrated in mid/late layers (Figure 3b). The information the decision needs is not in the field but distributed over the tokens that followed it at prefill (Figure 3c).

(3) Linear decodability. The field-conditioned conclusion is linearly decodable from those downstream tokens’ residual stream at prefill time—i.e. the model has already *computed and written down* the answer, not merely copied the field.

(4) Knockout and dose-response. Ablating the high-effect downstream tokens flips the decision back to stale, and the effect grows with the number/position of memoizing tokens, ruling out a diffuse explanation: specific aggregator/delimiter positions carry the conclusion. This mirrors the delimiter-token aggregation reported by Lindsey et al. [11] for forward planning; here the stored quantity is a backward-looking, field-conditioned *conclusion*.

What the note contains. If the note were a verbatim copy of the field, restating the value would suffice and emphatic wording would be inert. Instead, a wording ablation (Figure 3d) shows that once the corrected value is present, the override phrasing is *redundant* (bare value, tagged update, and explicit override all reach ≈ 1.0 P(safe)), while aggressive “disregard your earlier conclusion and re-evaluate” phrasing actively *hurts* (0.81). The note behaves like a committed conclusion that a late, salient correction can overwrite—but that confrontational instructions can destabilize. We name the phenomenon **attention-mediated memoized inference**. The remainder of this section stress-tests the account—further controls, more architecture families, and a component-level circuit—before the rest of the paper turns it into capabilities.

Separating conclusion from content—causal, temporal, and writable. Four further controls (three models each: Qwen3-8B/4B, Llama-3.1-8B; Appendix C) close the gap between “the conclusion is *decodable* downstream” and “the decision *uses* a memoized conclusion.” **(i) Disso-ciation.** We hold the field value byte-identical and flip a single rule token (a polarity *trigger*) so the *conclusion* inverts while the *content* does not; transplanting the downstream notes alone then carries the whole flipped conclusion (recovery 0.998–1.009) while patching the differing rule token carries none (-0.007 to $+0.007$)—the decision cannot be re-encoding field content, since the content is constant. (A linear probe finds *both* conclusion and field identity decodable downstream, so decodability alone cannot separate them: the causal patch is the necessary instrument.) **(ii) Timing.** Within the single prefill pass, the conclusion becomes linearly decodable on the downstream aggregator at depth 0.31–0.39, roughly twelve layers *before* the decision token’s logit-lens commit (depth 0.73–0.77): the note is written at prefill, before it is read. **(iii) Specificity.** Transplanting the eight highest-effect downstream positions recovers 0.74–0.79 of the decision; eight *random* downstream positions recover ≤ 0.035 —a few specific aggregator tokens, not a diffuse code. **(iv) Writability.** Injecting a *false* note (the opposite conclusion’s downstream KV) over an otherwise-consistent cache makes the decision follow the written note *against its own live field* (recovery ≈ 1.0); a handful of note tokens suffice. The account also holds off the synthetic template—field-only recovery stays




Figure 3: **Four causal probes for memoized inference.** (a) Refreshing the field’s own KV recovers ≈ 0 of the decision; recomputing the downstream recovers it fully. (b) Recovery is suffix-concentrated, accruing only as many post-field tokens are patched. (c) The decision reads almost entirely from downstream notes, not the field token. (d) Once the value is present, override wording is redundant and “re-evaluate” phrasing hurts—the note is a committed conclusion.

≈ 0 for multi-hop reasoning and free-form conversational phrasing, bounded only for near-verbatim attribute lookup, where the field is partly a copy (Appendix C).

Not a model-family artifact: four families. To rule out the aggregator-token account being a Qwen3/Llama tokenizer artifact, we replicate five probes (the locality probe and the four deep controls) on two further architecture families, **Gemma-2-9B** and **Mistral-7B** (a tokenizer-robust readout; for Gemma-2 we keep its attention/logit soft-capping intact). Every result holds: field-only recovery 0.005/0.137 vs. full-downstream 1.0; the conclusion/content dissociation (trigger-only ≈ 0 with the field held identical vs. notes ≈ 1.0); top-8 vs. random-8 specificity (0.95 vs. 0.48); false-note injection (0.98–1.0, follow-rate 1.0); and write-before-read timing (write depth 0.19–0.26 vs. decision commit 0.47–0.48). The mechanism is consistent across *four* families (Qwen3, Llama-3.1, Gemma-2, Mistral; Appendix C).


A component-level circuit. The probes above localize the memoized conclusion; a further five interventions (Appendix D, replicated across four families: Llama-3.1, Qwen3, Gemma-2, Mistral) resolve it to *components*. The picture is a *distributed write, concentrated read*: at mid layers *attention* writes the field-conditioned conclusion onto aggregator tokens (≈ 0.6 of the write on Llama), in a redundant code that is sparsely *decodable*—a trained SAE feature reaches AUC 1.0—yet causally spread over ~ 10 – 30 features and a shared low-rank direction ($\approx 25\times$ a random direction); a small, nameable set of late *read heads* (e.g. those attending decision $\rightarrow$ aggregator) then funnels it into the decision logit (12 heads recover 0.78 of the decision, random heads ≈ 0); and causal scrubbing confirms the note alone governs the decision (resampling it to the opposite conclusion flips the decision; resampling everything else does not). The feature that *decodes* the conclusion is not causally sufficient alone—the same decodability-vs-causation split we control for above, now at the level of single features. With the mechanism established, the rest of the paper shows it makes the cache both editable (Section 4) and composable (Section 5).

4 Consequence I: the cache is editable

Figure 4 contrasts the two cache operations this section and the next make precise: editing appends a salient correction; composing repositions and splices precompiled KV.

Naive editing fails—by design. The mechanism predicts the in-place edit will fail, and it does: refreshing the field’s KV while reusing the stale downstream leaves the decision at the old value

(a) **Editable:** append a salient erratum ($O(1)$); reuse the whole prefix



(b) **Composable:** reposition + splice precompiled notes ($O(L)$); skip repreload




Figure 4: **Two cache operations.** (a) *Editable*: append a salient erratum that amends the stale notes ($O(1)$ extra work, the whole prefix is reused). (b) *Composable*: precompute a skill’s KV in isolation, RoPE-reposition it to the target positions, and splice it in ($O(L)$, no repreload).

(Figure 1b). The prefix before the field is genuinely reusable (deviation 0.0); the problem is that the conclusion was already memoized after it.

Two ways to fix it, neither a full repreload. If the stale notes carry an old conclusion, there are two interventions. (i) *Recompute the affected notes* after the field. Recomputing the *whole* affected suffix is reliable—it is exactly *full-downstream* (recovery 1.0)—but pays the post-field prefill; re-computing only the top- K highest-effect notes (*field+selective@K*) is cheaper but *unreliable*, because the minimal K is strongly model-dependent (below). (ii) *Append an erratum*: one salient line late in the context—“[STATE UPDATE] field $\rightarrow$ new; overrides any earlier value and conclusion”—so the decision token attends to a fresh, authoritative note (verbatim template in Appendix B). The erratum is the cheap, robust default: appended after the field (*field+erratum*) it matches the strong *hoist-to-end* oracle *without rewriting the prompt* ($P(\text{correct})$ 1.00 on the gated-task frontier, Figure 2a; Table 6), and being append-only it composes with prefix caching (Section 9).

Chain-of-thought, not model size, gates the cheapest edit. A third, near-free option—refresh the field’s own KV alone ($\sim 1\%$ compute) and nothing else—works only when some later computation actually re-reads the field. A reasoning *chain* does exactly that, so with chain-of-thought the field-only edit alone recovers the decision (1.00 on Qwen3-8B); without the chain the *identical* edit on the *same* model is ignored (0.00) and the decision commits to the stale note (Figure 2c). The divider is therefore the CoT *mode*, not raw scale—reasoning-native models default to CoT, instruction-tuned ones to direct answers—so without CoT one of the two real fixes above is required. The mechanics are that simple, and we keep them brief here; the per-model stickiness, the *field+selective@K* sweep ($K^* \approx 4$ at 8B but > 64 at 4B, varying non-monotonically with scale; Figure 5a,b), and the layer-wise account are detailed in Appendices C and E. We report *field+selective@K* honestly as a genuine but *unreliable* surgical tool, not a default.

A frontier, not a winner. A controlled comparison (full table in Appendix E) shows no single dominant method. Hoist-to-end is cheapest but demands prompt surgery and pre-identification of every field; *field+erratum* matches it with no surgery at a one-line append; *in_place* is near-free but only with chain-of-thought; a KV-deviation-ranked selective recompute (CacheBlend-style)

Table 1: **KV editing vs. weight editing** for mutable per-request state (Llama-3.1-8B). Weight edits succeed at the target yet are global (contaminate concurrent requests), damage unrelated decisions, and are 30–50 $\times$ slower per edit.

method	flips decision?	edit latency	cross-req. contamination	collateral
field+erratum (KV)	✓	114 ms	0	0
in_place (KV)	✗ (no CoT)	71 ms	0	0
ROME (rank-one)	✓	5.6 s + 11.6 s cov	1.0	0.5
LoRA fine-tune	✓	3.1 s	1.0	0.5




Figure 5: **Editing is model-dependent** (the editing landscape and the chain-of-thought split are previewed in Figure 2). (a) The stickiness of the memoized conclusion is model-dependent: the near-free field-only recovery under chain-of-thought varies non-monotonically across model sizes. (b) field+selective@K: the minimal recompute K^* to reach full quality is likewise model-dependent.

underperforms here because it chases changed keys rather than the tokens that *memoized the conclusion*. The practical recommendation is field+erratum as the robust default, with in_place as a free fast-path under chain-of-thought.

Why not edit the weights? A natural objection: to act on a changed field, why not edit the model’s *weights* (ROME/MEMIT) or fine-tune, rather than the cache? We compare on the paper’s gated task (status pending→shipped, so *cancel*→*deny*) against a faithful rank-one ROME—validated on the canonical factual edit first (“the Eiffel Tower is in” *Paris*→*Rome*, locality intact), so the baseline is not crippled—and a LoRA fine-tune (Table 1; methodology in Appendix I). All of field+erratum, ROME, and LoRA *succeed* at flipping the target decision. But a weight edit is *global*: the same model instance can no longer hold status=shipped for one request and pending for another, so *all* concurrent orders that are genuinely still pending are wrongly flipped (cross-request contamination 1.0), and half of an unrelated decision battery drifts (0.5)—the ROME/fine-tune specificity tax—at 3–6 s per edit (plus a one-time covariance pass for ROME). The append-only erratum lives in a *per-sequence* cache: zero cross-request contamination, zero collateral, 114 ms, and it composes with prefix caching (Section 9). Weight editing targets *durable, global facts*; mutable per-request, per-turn state is the wrong job for it—which is precisely the niche the editable KV cache fills.

5 Consequence II: the cache is composable

From mechanism to prediction. If a region’s notes are localized and re-derivable from context the decision can still see, then the notes should be *position-portable*: we can compute a reusable chunk’s KV once, in isolation, move it to a new absolute position, and splice it in. Concretely we precompile a long *skill* (a policy or tool specification), and because the attention library caches post-RoPE keys, we re-rotate the chunk’s keys from their source positions to the target positions (values are position-free) before concatenating (Figure 4b). This is the position-aware reuse of Liu et al. [12]; our point is that the *mechanism predicts it should preserve behavior*.

Transplant is behaviorally indistinguishable from full recompute. It does. The spliced skill matches a full refill in next-token logits with cosine 0.90–0.999 across the full model family—Qwen3-1.7B through 32B (including FP8 and the 30B-A3B Mixture-of-Experts), Gemma-2/3, Mistral-7B, Llama-3.1-8B and 70B, and DeepSeek-R1-Distill-Llama-8B (Figure 6b)—and on the models competent at the task it preserves *correct* skill-following across 8 diverse domains and 3 families (24/24, cosine 0.98–0.999), including 16/16 under chain-of-thought.

Context-robustness and the seam. A chunk precompiled in isolation matches one that attended to the real preceding context, because the decision re-derives from context it still sees. The one residual error is a *seam* at the chunk’s start—the first tokens that, in a full prefill, would have attended to the now-missing prefix. Recomputing a few boundary tokens (*seam-repair*) closes it; this is exactly the boundary recompute of CacheBlend/EPIC/MPIC, and Section 3 explains why it is the boundary, specifically, that needs repair.

Linear-time TTFT. Full refill of a length- L skill is $O(L^2)$; transplant is $O(L)$ (a re-rotation pass plus the suffix). Time-to-first-token speedups grow with skill length: $3\times$ at 2k tokens, $9.8\times$ at 8k, and $13.9\times$ at 32k on an 8B model (Figure 6a). A library of skills composes (decisions preserved for $N = 1\text{--}4$ concurrent skills).

Generality of content, placement, and agency. Transplantation is not specific to rule-like skills. It preserves decisions for *facts/RAG* passages as well as rules; for chunks inserted in the system area *and* mid-trajectory as tool results; and—measured with actual function calls rather than a proxy—it preserves *agentic tool-calling* ($N=108$ with bootstrap CIs: function-call accuracy 1.00 on Mistral-7B, Llama-3.1-8B/70B, and Qwen3-8B; 0.97 [0.94, 1.0] on Qwen3-32B-FP8 and 30B-A3B, whose transplanted-vs-full tool-call *agreement* is 1.00). The one consistent exception is sliding-window attention (Gemma), which we diagnose and fix in Section 8. Full per-domain scorecards are in Appendix F.

6 One substrate: the keystone and a unified agent

The keystone: editing inside a transplant. If editing and composing are truly two operations on the same object, then editing a field that lives *inside a transplanted skill* should behave exactly as editing a field in a normally-prefilled context. We test this directly: transplant a skill whose body contains a mutable field, then apply each editing method to that field and measure recovery, comparing the *composed* cache (skill spliced in) against a fully *recomputed* cache. The editing mechanism reproduces verbatim (Figure 7a): the in-place edit is weak (≈ 0.05), selective recompute recovers ($\text{sel}@32 \approx 0.80$), the erratum is strongest, and crucially *composed* $\approx$ *recomputed* for every method (points lie on the diagonal) across Gemma-2-9B and Llama-3.1-8B. The notes a transplant pastes in are the same notes an edit amends—one substrate.

A unified edit+compose agent. We embody both operations in a single live agent loop: a long policy is *composed* once and never re-prefilled; as the world changes across turns the mutable state




Figure 6: **Composing the cache.** (a) Pasting a precompiled skill is $O(L)$ vs. full refill’s $O(L^2)$; TTFT speedup reaches $13.9\times$ at 32k tokens. (b) The transplanted skill matches full recompute in next-token logits (cosine 0.90–0.999) across the model family.




Figure 7: **Edit and compose are one substrate.** (a) Keystone: editing a field *inside* a transplanted skill reproduces the editing mechanism, and the *composed* cache matches the *recomputed* cache for every method (points on the diagonal). (b) A unified edit+compose agent over thirteen models (10 domains $\times$ 10 instances, 300 decisions; 120 on the two Gemma models): unified-vs-full agreement (bars) and cumulative-TTFT speedup (markers).

is *edited* by appended errata; and each turn reuses the longest cached prefix and prefills only the delta. Against a refill-every-turn baseline, over 10 agent domains $\times$ 10 instances (300 decisions) per model (40 instances, 120 decisions, on the two Gemma models) and across thirteen models (the twelve transplant models of Figure 6 plus Qwen3-0.6B; Appendix A), the unified path is decision-identical to full recompute with agreement 0.81–1.00 (e.g. Llama-3.1-8B 0.963, Mistral-7B 0.983, Llama-3.1-70B 1.00) at cumulative-TTFT speedups up to $14.9\times$ (Figure 7b); the speedup scales with policy length $\times$ turns. This is editing and composing operating together, losslessly and faster, across the model family. The leave-stale+erratum cache also does *not* compound error over long trajectories: across a 28-turn stress test where the gating field toggles every turn, the decision logits stay faithful to a full refill (cosine 0.99+, flat with trajectory length) with no systematic drift (Appendix G).

7 Application: editable and composable user memory

A concrete, high-value instance of the edit+compose substrate is *user memory*: the large, dynamically summarized profile of facts an assistant re-reads every turn. Memory is big (thousands–tens




Figure 8: **E1 — placement and the pre-digestion cost (Qwen3-4B, CoT).** Reading memory *late* carries a small, statistically significant accuracy cost vs. *early* that *grows with memory length*; late placement is still preferred because it enables $O(L)$ editing/transplant (efficiency quantified in E5).

of thousands of tokens), reused across turns, and *mutated mid-session* by tool calls—so where it lives in the prompt is a dilemma the mechanism of Section 3 explains. Placed at the *front* (`[sys][MEM][traj]`), the trajectory memoizes memory-conditioned conclusions downstream, so a memory change forces a costly downstream reprefill (and an in-place edit is ignored). Placed at the *end* (`[sys][traj][MEM]`), memory’s KV depends on the whole trajectory and must be re-attended every turn. We resolve this by treating memory as a *skill that is also edited*: precompile it once in isolation, place it late (`[sys][traj][MEM][query]` so the decode reads it directly), RoPE-reposition it each turn, repair one boundary token, and edit it in place when it changes.

Placement and the pre-digestion cost (E1; Figure 8). Pre-digestion is real, and we quantify it. Under full recompute and chain-of-thought, reading memory *late* carries a small but *statistically significant* accuracy cost vs. *early* (Qwen3-4B, 192 personas per cell: GEE-logistic $\beta_{\text{late}} = -0.44$, $p = 0.018$), and the cost *grows with memory length*: the early–late gap is ≈ 0 at 2k tokens but $+0.09$ at 16k and $+0.16$ at 32k. The mechanism explains it—late placement forgoes prefill-time pre-digestion and relies on the decode/CoT to integrate raw memory. This is a *tradeoff, not free*: late placement is still preferred because it enables $O(L)$ editing/transplant at $2.3\text{--}4.3\times$ lower TTFT and the transplant is faithful at a fixed placement (both shown below), so the net is small accuracy for large efficiency—with early placement preferable when memory is very long and accuracy-critical. (Direct one-shot decisions are at chance for the reasoning-native Qwen3 family to 32B, breaking only at Llama-3.1-70B, direct 0.81; CoT is the operative regime across scale—see Appendix H.)

Memory transplant is faithful, to 70B (E2; Figure 9). Across ten models from 0.6B to 70B, a precompiled+repositioned memory chunk reproduces the full-recompute decision logits with cosine 0.94–0.9996; a *single* seam-repair token closes the start-of-chunk boundary gap (Llama-3.1-8B late: cos 0.94 $\rightarrow$ 0.994). The cleanest decision-governance test is Llama-3.1-70B, whose decisions *genuinely vary* (so agreement is non-trivial, unlike the constant-answer regime of smaller models): the late, seam-repaired transplant reproduces the full-recompute decision 0.93 of the time at logit cosine 0.997, and *late placement beats early* (0.93 vs. 0.83) exactly as the mechanism predicts (the decode reads memory directly rather than through pre-digested notes across the transplant boundary). The no-rotation control collapses (late decision agreement 0.18 vs. 0.78 rotated), confirming the re-rotation is necessary.




Figure 9: **E2 — memory transplant is faithful, to 70B.** (a) Decision agreement with full recompute under late placement; a single seam-repair token closes the start-of-chunk gap across ten models. (b) The re-rotation is necessary—the naive no-rotate control collapses, while rotated+1 seam token matches full recompute.

Memory is editable mid-session (E3; Figure 10). When a stored fact is toggled, reusing the *stale* memory recovers the flipped decision essentially never (≤ 0.03). Every real edit recovers it, and—consistent with Section 4—under chain-of-thought the *near-free in-place edit* (one token recomputed) suffices, and *strengthens with scale*: correctness $0.92 \rightarrow 0.99 \rightarrow 1.00$ across Qwen3-1.7B/4B/14B (and 0.98 at 32B). On models where the chain does not re-read the field, the append-only ERRATUM is the robust fallback (McNemar $\text{ERRATUM} > \text{in_place}$ on Llama-3.1-8B, $p = 0.031$). This is the editing axis that concurrent KV-cache memory systems lack: MemArt [1] and EPIC [6] retrieve and splice *static* memory blocks position-independently but have no in-place memory *update*; our additions are the editing operation, the decision-governance lens, and the mechanism (Section 3) that explains why boundary recompute suffices.

Keystone: editing inside transplanted memory reproduces the mechanism. Editing a field *inside a transplanted memory chunk* in direct mode (Llama-3.1-70B, whose decisions vary so recovery is measurable) reproduces Section 3’s memoization verbatim: refreshing the field’s KV alone recovers only 0.55 of the flipped decision (the conclusion was memoized *downstream*, not in the field), and recovery climbs monotonically as more downstream tokens are recomputed ($0.55 \rightarrow 0.84$ at $K=16 \rightarrow 0.94$ full), exactly as for a normally-prefilled context. Under chain-of-thought this stickiness *dissolves*—a selective@ K sweep is flat at ≈ 0.98 for all K (the chain re-reads the field, so $K^* \approx 1$). Edit and compose thus act on one substrate, and the direct/CoT dissociation matches the editing law of Section 4.

Real memory: LoCoMo external validity (Figure 11). The synthetic gated decisions isolate the mechanism; to test real memory we run the long-conversation QA benchmark LoCoMo [15] (MemArt’s setting): the multi-session dialogue (median $\sim 19.7k$ tokens) is the memory, precompiled and spliced before each question. Over *all* 1,540 answerable questions per model, transplant is *statistically equivalent* to full recompute in QA accuracy on Qwen3-4B, 14B, and 32B (TOST, margin 0.03: $|\Delta| \leq 0.015$) and within a small -2.7 points on Llama-3.1-8B, with answer-token logit cosine $0.991\text{--}0.998$ throughout. The *compose* axis thus holds on real conversational memory, not only synthetic decisions.




Figure 10: **E3 — memory is editable mid-session.** Reusing stale memory recovers a toggled fact essentially never; every real edit recovers it. Consistent with Section 4, the near-free in-place edit suffices under chain-of-thought and *strengthens with scale*, with the append-only erratum as the robust fallback.

Granularity is a free knob (E4; Figure 12). Splitting memory into S independently-precompiled blocks makes a localized edit $S \times$ cheaper (recompute one block) and is *decision-lossless* up to $S=16$ (agreement flat at 1.00 on Qwen3-4B), because the independent facts are integrated at read time, not within memory. Logit cosine degrades with S ($0.998 \rightarrow 0.953$) but *identically whether the gating facts are contiguous in one block or split across blocks*—splitting *independent* relevant facts across independently-precompiled blocks does not specifically hurt. Genuinely cross-referential facts, however, do: in a controlled test (Appendix G) where a decision needs a two-hop chain (a DEFINITION line names which setting gates, whose value lives elsewhere), splitting the linked pair across a block boundary drops decision agreement with full recompute to 0.46 vs. 0.76 when the pair stays in one block—a 0.30 penalty (Llama-3.1-8B, $n=80$, full-recompute accuracy 0.85; McNemar $p < 10^{-6}$), because block B ’s isolated precompute never attended to its referent A . The same split costs an independent fact pair only 0.04. Practical guidance: keep cross-referential facts in one block.

End-to-end agent (E5; Figure 13). We implement a live agent (Appendix H) that composes memory once, re-rotates it each turn, and edits it on tool-driven changes. Over 12-turn sessions (16–120 sessions per model) with $\approx 2k$ -token memories, against a refill-every-turn-at-the-end baseline it cuts cumulative time-to-first-token by $2.3\text{--}4.3 \times$ (growing with model size, to 32B), and against front-placement refill-on-change by $1.4\text{--}3.3 \times$, while reproducing the full-refill next-token logits at a token-matched oracle (cosine 0.97–0.99). One honest caveat: greedy chains-of-thought are sensitive to sub-percent logit differences, so the exact reasoning *chain* is not always reproduced (chain agreement 0.31–0.78) even though the next-token decision is faithful and CoT *accuracy* is comparable; the clean decision-governance equivalence is the short-context editing result above (Section 4-style, agreement 0.89–0.95). User memory is thus a second instance of one substrate that is both editable and composable.

8 Reach and the 2026 frontier

Scale, quantization, MoE. The transplant is faithful from 0.6B to 32B, on FP8 checkpoints, on a 30B-A3B Mixture-of-Experts, and on a 4-bit 70B model (feasibility 8/8, logit cosine 0.986); the unified agent of Section 6 likewise spans all thirteen.




Figure 11: **LoCoMo — external validity on real conversations.** Over all 1,540 answerable questions per model, transplanting the multi-session dialogue memory is statistically equivalent to full recompute in QA accuracy (TOST) on the Qwen3 models and within 2.7 points on Llama-3.1-8B.

Multimodal: images take notes too. In an agent trajectory an image costs a full prefill—the vision tower *plus* prefilling its $>1k$ soft-tokens through the LM. Because image notes are also position-portable, we cache an image’s LM-side KV once and splice it, re-running only text. Across 120 diverse VQA tasks per model (perception/reasoning/agentive), the spliced image is near-lossless versus full re-encode—agreement 0.958–1.0 on Qwen2.5-VL-3B/7B/32B and Qwen3-VL-8B (Figure 14b)—and reusing a cached image is $2.4\text{--}8.4\times$ faster TTFT. Moving an image to a different trajectory position requires re-rotating only the temporal axis of M-RoPE; we handle both the *sectioned* (Qwen2.5-VL) and *interleaved* (Qwen3-VL) layouts, with position-shifted transplant lossless (agreement 0.99, $\Delta=161$ positions). This subsumes MPIC’s multimodal reuse [31], replacing its boundary-recompute with exact re-rotation.

The substrate is any per-token attention KV. The operations depend on the cache *representation*, not the attention kernel, so we map exactly where they hold (Figure 14a):

- **Free**—throughput optimizations that keep per-token KV: FlashAttention, paged attention/vLLM, and GQA/MQA (already used by every model we ran).
- **Adapter (implemented, validated)**—representation changes (diagrammed in Figure 23). *MLA* (DeepSeek-V2/Coder-V2) caches a position-free latent plus a small decoupled-RoPE sub-vector; our decoupled k_{pe} reposition re-rotates only that sub-vector, giving logit cosine 0.98 and composed-vs-full agreement 1.00 on DeepSeek-Coder-V2-Lite. *Interleaved M-RoPE* (above) is the other adapter.
- **Fixed**—sliding-window (Gemma): the default cache truncates sliding layers to the window, breaking uniform splices beyond it; keeping the *full* per-token KV and letting the attention *mask* enforce the window restores correctness (the previously-failing unified agent now runs at agreement 0.93–0.94; Figure 23c).
- **Partial**—hybrids with full-attention layers (Falcon-H1): the attention KV is transplantable but the per-layer Mamba scan-state is recurrent, not per-token, so a correct transplant must re-scan the Mamba path and saves only the attention fraction.
- **Open frontier**—sequence-dimension KV compression (DeepSeek-V4’s CSA/HCA) merges tokens into sub-token-count entries, so edit/splice become block-granular; DeepSeek-V3.2’s sparse




Figure 12: **E4 — granularity is a free knob.** Splitting memory into S independently-precompiled blocks makes a localized edit $S\times$ cheaper and stays decision-lossless to $S=16$; only genuinely cross-referential facts must be kept in one block (Appendix G).

attention (DSA) is MLA plus top- k selection and inherits our MLA adapter.

- **Out of scope**—no per-token attention KV: pure-recurrent (RWKV), pure-SSM (Mamba), and diffusion LMs. The prompt-level erratum still applies there, but it is not a KV operation.

9 Systems payoff

A real agentic environment. On the τ^2 -bench retail environment [2]—single tool-decisions and a multi-turn autonomous-agent loop scored by the environment’s own tool enforcement—an agent that reuses a stale cache after a state change collapses, while `field+erratum` preserves task success at a fraction of the recompute. On the real ~ 1.4 k-token retail policy, transplant reproduces the clean decision (composed = full on Llama-3.1 and Mistral); the one hard case, a long buried field that must *flip* a conclusion, requires the robust `field+erratum` edit rather than transplant-plus-bare-erratum—the same long-context lesson the editing axis predicts, now in a real environment.

A comprehensive online serving benchmark. Because the erratum is append-only, it composes with standard automatic prefix caching (APC): the static prefix stays cache-aligned and only the short erratum (and the decode) is new work, whereas writing the new value *into* the prefix changes a cached block’s content hash and invalidates every downstream block. We test this on vLLM’s V1 engine as a real online server—AsyncLLMEngine with CUDA graphs, continuous batching, APC, and *Poisson* request arrivals at controlled offered load (not an offline batch)—over a shared ~ 8 k-token agent policy with one mutable field, measuring TTFT percentiles, throughput, and the engine’s own APC hit-rate (Figure 15a; full configuration in Appendix I). The append-only edit keeps the prefix a cache hit (98.5% **vs.** 1.0% **APC hit-rate**); the in-prefix baseline is prefill-bound and *saturates at* ≈ 1.5 req/s, so its p_{90} time-to-first-token collapses under load (22–55 s) while




Figure 13: **E5 — end-to-end agent.** (a) Per-decision median TTFT for the proposed compose+edit agent vs. oracle, front, and end-placement baselines. (b) Cumulative TTFT speedup of 2.3–4.3× over refill-every-turn, growing with model size, at faithful next-token decisions.

the erratum stays 86 ms–1 s (53–398× lower TTFT). The throughput advantage *grows with offered load*—1.6× at 2 req/s up to **14.5×** at saturation—exactly as predicted for a compute-bound vs. cache-bound regime. The image-KV reuse of Section 8 contributes a complementary serving win—2.4–8.4× faster first-token as image size grows (Figure 15b)—by skipping the vision tower and image-token prefill entirely.

10 Limitations

Our operations require a per-token attention KV cache, so pure-recurrent, pure-SSM, and diffusion models are out of scope, and hybrid attention+SSM models are only partially served (the recurrent state is not transplantable; Section 8). The MLA and sliding-window adapters are validated but carry residual edge cases: a single transplanted *chunk that itself exceeds the sliding window* drops to logit cosine ≈ 0.89 (real skills are far smaller), and our small MLA checkpoints ship legacy-cache custom modeling that we shim. Sequence-dimension KV compression (DeepSeek-V4-class) and cross-attention image caches are open: the unit of edit/transplant there becomes a compressed block rather than a token, which we have analyzed but not implemented. The `field+selective@K` surgical edit is unreliable (model- and domain-dependent stickiness) and we present it as such, not as a default. Finally, while several studies use synthetic policies for controlled measurement, we mitigate this with the real τ^2 -bench retail policy and real images; broader real-workload evaluation remains future work. The caching machinery we build on is prior art (Section 2); we claim the mechanism, the unification, the decision-governance lens, and the attention-variant adapters.

11 Conclusion

A surgical edit to a field’s KV is ignored not because the cache is fragile but because the model already did the work: at prefill it memoizes the field-conditioned *conclusion* onto downstream aggregator tokens, and the decision reads those notes. Reframing the KV cache as a notebook of memoized conclusions makes two capabilities fall out of one mechanism—you can *edit* the notes (amend them with a late, salient erratum rather than recompute) and *compose* the notes (reposition and splice a precompiled skill, in $O(L)$ time)—and a keystone experiment shows they act on a single substrate. The substrate is any per-token attention KV cache: it holds across scale, quantization, Mixture-of-Experts, and multimodal image caches, transfers freely across throughput optimizations,

(a) the substrate = any per-token attention KV

attention variant	status
FlashAttn / paged / vLLM	free
GQA / MQA	free
MLA (DeepSeek-V2/V3)	adapter
M-RoPE (sectioned/interleaved)	adapter
Sliding-window (Gemma)	fixed
Hybrid attn+SSM (Falcon-H1)	partial
Seq.-dim compress (V4 CSA/HCA)	open
RWKV / Mamba / diffusion	out

(b) images are position-portable too




Figure 14: **Reach.** (a) The substrate is any per-token attention KV cache; we map each attention variant as free / adapter / fixed / partial / open / out-of-scope. (b) Image-KV transplant is near-lossless across vision-language models—images are position-portable too.

and reaches Multi-head Latent Attention, sliding-window, and interleaved-M-RoPE models through small adapters—up to the edge of the 2026 sparse/compressed-attention frontier, where the unit of editing and composition becomes a block rather than a token. In practice this turns two everyday agent costs—reloading a library of skills and re-reading a mutating user-memory document—into $O(L)$ cache operations that preserve the decision, with up to $14.5\times$ online serving throughput. We hope the “models take notes at prefill” lens is useful beyond editing and composition: the cache is a readable, writable record of what a model has already concluded.

Acknowledgements

We thank Junhao Hu, first author of EPIC (Efficient Position-Independent Caching) [6] and a co-author of CacheSlide [12], for discussions on editable KV cache that inspired this research project.

References

- [1] Anonymous. MemArt: KVCache-centric memory for LLM agents, 2025. Under review (Open-Review), posted October 2025.
- [2] Victor Barres, Honghua Dong, Soham Ray, Xujie Si, and Karthik Narasimhan. τ^2 -bench: Evaluating conversational agents in a dual-control environment. *arXiv preprint arXiv:2506.07982*, 2025.
- [3] DeepSeek-AI. DeepSeek-V2: A strong, economical, and efficient mixture-of-experts language model. *arXiv preprint arXiv:2405.04434*, 2024.
- [4] DeepSeek-AI. DeepSeek-V3 technical report. *arXiv preprint arXiv:2412.19437*, 2024.
- [5] In Gim, Guojun Chen, Seung-seob Lee, Nikhil Sarda, Anurag Khandelwal, and Lin Zhong.




Figure 15: **Systems payoff.** (a) Online vLLM serving (V1 engine, CUDA graphs, continuous batching, APC, Poisson arrivals): the append-only erratum keeps the prefix cache-aligned (98.5% vs. 1% APC hit-rate), so its throughput advantage grows with offered load—up to $14.5\times$ at saturation—while p_{90} TTFT is 53–398 $\times$ lower. (b) Reusing a cached image (skipping the vision tower and image-token prefill) accelerates time-to-first-token, more so for larger images.

Prompt cache: Modular attention reuse for low-latency inference. In *Proceedings of Machine Learning and Systems (MLSys)*, 2024.

- [6] Junhao Hu, Wenrui Huang, Weidong Wang, Haoyi Wang, Tiancheng Hu, Qin Zhang, Hao Feng, Xusheng Chen, Yizhou Shan, and Tao Xie. EPIC: Efficient position-independent caching for serving large language models. In *Proceedings of the 42nd International Conference on Machine Learning (ICML)*, 2025.
- [7] Gabriel Ilharco, Marco Tulio Ribeiro, Mitchell Wortsman, Ludwig Schmidt, et al. Editing models with task arithmetic. In *International Conference on Learning Representations (ICLR)*, 2023.
- [8] Chao Jin, Zili Zhang, Xuanlin Jiang, Fangyue Liu, et al. RAGCache: Efficient knowledge caching for retrieval-augmented generation. *arXiv preprint arXiv:2404.12457*, 2024.
- [9] Kenneth Li, Oam Patel, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. Inference-time intervention: Eliciting truthful answers from a language model. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [10] Yuhong Li, Yingbing Huang, Bowen Yang, Bharat Venkitesh, et al. SnapKV: LLM knows what you are looking for before generation. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- [11] Jack Lindsey et al. On the biology of a large language model. *Transformer Circuits Thread, Anthropic*, 2025.
- [12] Yang Liu, Yunfei Gu, Liqiang Zhang, Chentao Wu, Guangtao Xue, Jie Li, Minyi Guo, Junhao Hu, and Jie Meng. CacheSlide: Unlocking cross position-aware KV cache reuse for accelerating LLM serving. In *Proceedings of the 24th USENIX Conference on File and Storage Technologies (FAST)*, 2026.
- [13] Yuhan Liu, Hanchen Li, Yihua Cheng, Siddhant Ray, et al. CacheGen: KV cache compression and streaming for fast large language model serving. In *Proceedings of ACM SIGCOMM*, 2024.

- [14] Zichang Liu, Aditya Desai, Fangshuo Liao, Weitao Wang, et al. Scissorhands: Exploiting the persistence of importance hypothesis for LLM KV cache compression at test time. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [15] Adyasha Maharana, Dong-Ho Lee, Sergey Tulyakov, Mohit Bansal, Francesco Barbieri, and Yuwei Fang. Evaluating very long-term conversational memory of LLM agents. In *Proceedings of the Association for Computational Linguistics (ACL)*, 2024.
- [16] Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual associations in GPT. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- [17] Kevin Meng, Arnab Sen Sharma, Alex Andonian, Yonatan Belinkov, and David Bau. Mass-editing memory in a transformer. In *International Conference on Learning Representations (ICLR)*, 2023.
- [18] Bowen Peng, Jeffrey Quesnelle, Honglu Fan, and Enrico Shippole. YaRN: Efficient context window extension of large language models. In *International Conference on Learning Representations (ICLR)*, 2024.
- [19] Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, et al. Toolformer: Language models can teach themselves to use tools. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [20] Jianlin Su, Murtadha Ahmed, Yu Lu, Shengfeng Pan, Wen Bo, and Yunfeng Liu. RoFormer: Enhanced transformer with rotary position embedding. *Neurocomputing*, 568, 2024.
- [21] Jiaming Tang, Yilong Zhao, Kan Zhu, Guangxuan Xiao, Baris Kasikci, and Song Han. Quest: Query-aware sparsity for efficient long-context LLM inference. In *International Conference on Machine Learning (ICML)*, 2024.
- [22] Eric Todd, Millicent Li, Arnab Sen Sharma, Aaron Mueller, Byron C Wallace, and David Bau. Function vectors in large language models. In *International Conference on Learning Representations (ICLR)*, 2024.
- [23] Jesse Vig, Sebastian Gehrmann, Yonatan Belinkov, et al. Investigating gender bias in language models using causal mediation analysis. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- [24] Kevin Wang, Alexandre Variengien, Arthur Conmy, Buck Shlegeris, and Jacob Steinhardt. Interpretability in the wild: a circuit for indirect object identification in GPT-2 small. In *International Conference on Learning Representations (ICLR)*, 2023.
- [25] Guangxuan Xiao, Yuandong Tian, Beidi Chen, Song Han, and Mike Lewis. Efficient streaming language models with attention sinks. In *International Conference on Learning Representations (ICLR)*, 2024.
- [26] Jingbo Yang et al. KVLink: Accelerating large language models via efficient KV cache reuse. *arXiv preprint arXiv:2502.16002*, 2025.
- [27] Jiayi Yao, Hanchen Li, Yuhan Liu, Siddhant Ray, Yihua Cheng, Qizheng Zhang, Kuntai Du, Shan Lu, and Junchen Jiang. CacheBlend: Fast large language model serving for RAG with

- cached knowledge fusion. In *Proceedings of the European Conference on Computer Systems (EuroSys)*, 2025.
- [28] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, et al. ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations (ICLR)*, 2023.
 - [29] Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. τ -bench: A benchmark for tool-agent-user interaction in real-world domains. *arXiv preprint arXiv:2406.12045*, 2024.
 - [30] Zhenyu Zhang, Ying Sheng, Tianyi Zhou, Tianlong Chen, Lianmin Zheng, et al. H2O: Heavy-hitter oracle for efficient generative inference of large language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
 - [31] Shiju Zhao, Junhao Hu, Rongxiao Huang, Jiaqi Zheng, and Guihai Chen. MPIC: Position-independent multimodal context caching system for efficient MLLM serving. *arXiv preprint arXiv:2502.01960*, 2025.

A Models evaluated

Table 2 lists the models used across the paper. All runs are on a single RTX PRO 6000 (Blackwell, 96 GB); FP8 and 4-bit checkpoints are the official quantized releases.

Table 2: Model zoo. “role” indicates the experiments a model appears in.

model	family / attention	precision	role
Qwen3-0.6B/1.7B/4B/8B/14B	Qwen3 (GQA)	bf16	mechanism, edit, compose, agent
Qwen3-32B	Qwen3 (GQA)	FP8	compose, agent, scale
Qwen3-30B-A3B	Qwen3 MoE (GQA)	bf16	compose, agent, MoE
Llama-3.1-8B / 70B	Llama (GQA)	bf16 / 4-bit	mechanism, edit, compose, agent
Mistral-7B	Mistral (GQA)	bf16	mechanism, compose, agent
Gemma-2-9B / Gemma-3-27B	Gemma (sliding-window)	bf16	mechanism, compose, agent, sliding-window fix
DeepSeek-R1-Distill-Llama-8B	Llama (GQA), reasoning	bf16	edit, compose, agent
DeepSeek-V2-Lite / Coder-V2-Lite	DeepSeek (MLA)	bf16	MLA adapter
Falcon-H1, Falcon-Mamba	hybrid / pure SSM	bf16	architecture boundary
Qwen2.5-VL-3B/7B/32B, Qwen3-VL-8B	VL (M-RoPE)	bf16	multimodal image-KV

B Worked examples: prompts, errata, and responses

The gated decision (the paper’s running example). The mechanism and editing experiments use prompts of the following shape (scenario `account_role`, the verbatim template from the released harness; the thirty neutral filler rules are abridged and long lines re-wrapped for display). The prompt is wrapped in each model’s chat template and the decision is read at the final position:

You are a customer-support agent for an online retailer.

SESSION CONTEXT

```
account_role: verified_admin    <-- the mutable FIELD (new value: suspended_user)
channel: web
agent_id: agent-7
```

POLICY DOCUMENT (read carefully; rules are binding)

```
ACCESS RULE (binding): If account_role is verified_admin, perform requested write
actions normally. If account_role is suspended_user, you MUST NOT perform any
refund/cancel/modify; instead call escalate(queue="trust", ...) and nothing else.
R1. Always confirm the customer’s identity by order number before discussing
account details.
[... R2-R30: twenty-nine more neutral filler rules ...]
```

AVAILABLE TOOLS

```
- lookup_order(order_id)          - issue_refund(order_id, amount, reason_code)
- cancel_order(order_id)          - share_payment_method(order_id)
- expedite_shipping(order_id)     - escalate(queue, summary)
- refuse(reason)
```

CONVERSATION SO FAR

```
user: Please refund 40.00 to my order A4471 for the shipping delay.
assistant: Let me check the account and policy before acting.
```

TASK

```
Decide the single next tool call. Respond with exactly one line:
tool_call: <name>(<args>)
Next action:
```

The four cache states of Section 3 are built from this prompt: *stale* prefills the old value, *oracle* the

new, *field-only* overwrites the KV of the `account_role` value span (two tokens of a ~ 958 -token templated prompt, 0.2% recompute), and *full-downstream* recomputes everything after it. In the position-resolved causal patch (Figure 3b,c), the highest-effect downstream positions for a representative instance sit not on the field span but on aggregator/delimiter tokens—the end-of-rule punctuation (around “...nothing else.”) and the TASK-block instruction tokens (around “tool_call: <”) that the decision attends through.

The erratum, verbatim. The editing fix appends one salient line ahead of the decision prompt (the `field+erratum` arm of Section 4):

```
[STATE UPDATE] account_role has just changed to suspended_user; this overrides
any earlier value. Apply the current value.
```

The wording ablation of Figure 3d varies this line from a bare restatement of the value to the aggressive “disregard your earlier conclusion and re-evaluate”, which underperforms.

Recorded responses (Qwen3-8B, reasoning mode; one released instance). From the record `thinking_qwen3_8b_think.json` (in `results/`), abridged to the answer head:

```
oracle (clean prefill of suspended_user), 249 thinking tokens:
  tool_call: escalate(queue="trust", summary="Request to refund $40.00...")
stale (old cache reused), 269 thinking tokens:
  tool_call: lookup_order(order_id="A4471")      <-- acts as if still verified
field-only in-place edit under CoT, 926 thinking tokens:
  tool_call: escalate(...)                      <-- chain re-reads the field
```

The in-place edit recovers the oracle decision under reasoning, but the recomputation has moved from prefill into the chain (926 vs. 249 thinking tokens, $\sim 3.7\times$); without reasoning, the same edit is simply ignored (Figure 3a).

The dissociation pair (Appendix C). The conclusion/content control holds the field byte-identical and flips one trigger token inside the rule:

```
ACCESS RULE (binding): If account_role is {trigger}, you MUST NOT perform any
refund, cancel, or modify action and must instead call escalate(queue="trust",
summary=...). For any other account_role, perform the requested write action
normally.
```

With the field fixed at `verified_admin`, setting `trigger=verified_admin` makes escalation correct, while `trigger=suspended_user` (field unchanged) makes the write action correct: the two prompts differ in exactly one token, the field is constant, and only the *conclusion* flips.

A transplanted skill (Section 5). A representative precompiled skill (the 20 neutral guidelines abridged):

```
# SKILL: REFUND_POLICY
You handle refund requests. Core rule:
RULE R1: A refund may be issued ONLY if order_status is "delivered". For any
other status (pending, shipped, cancelled, returned) you MUST refuse the refund
and escalate to a human.
- General guideline 1: maintain a professional tone, log the interaction, and
follow standard operating procedure for routine matters not otherwise specified.
[... general guidelines 2-20 ...]
End of REFUND_POLICY skill.
```

The skill is prefilled once in isolation (positions $0..L-1$), its keys re-rotated to the target offset, and spliced after the system prompt; the task suffix

```
Order #7731 has order_status = "pending". The customer requests a refund. Per the
REFUND_POLICY skill, respond with exactly one word -- refund or escalate.
Decision:
```

then reads the spliced notes, and the decision (`escalate`) matches a full reprefill (Figure 6).

C Deep mechanism controls: dissociation, timing, specificity, writability

These four controls (and an off-template generalization) tighten Section 3 from *decodability* to *causation, timing, and write-access*. Each runs on three models (Qwen3-8B, Qwen3-4B, Llama-3.1-8B); scripts `esys/mechd_*.py`, records `results/mechd_*`.

(i) **Dissociation (Table 4, left).** A polarity-parameterized rule names a single *trigger* value that selects the safe action, so flipping the trigger inverts the conclusion while the field value is byte-identical across the pair (the two prompts differ in exactly one token, inside the rule). Transplanting the post-trigger *notes* from the opposite-conclusion cache carries essentially the entire flipped conclusion, whereas patching the differing rule token carries none—with the field held constant, the decision is reading a memoized conclusion, not field content. A logit-probe finds both the conclusion and the field identity linearly decodable from the same downstream delimiter, so decodability cannot itself adjudicate; the causal transplant is required.

(ii) **Timing.** Using `output_hidden_states` on the prefill, the conclusion becomes linearly decodable (group-CV logistic probe, conclusion $\perp$ field by the 2×2 design) on the downstream aggregator at layer-depth 0.39/0.39/0.31 (Qwen3-8B/4B, Llama-3.1-8B), while the decision token’s logit-lens margin reaches its final sign only at depth 0.75/0.77/0.73—the note is written ~ 12 layers before it is read, within one prefill.

(iii) **Specificity (Table 4, right).** Ranking downstream positions by individual transplant effect, the top-8 recover 0.74–0.79 of the decision; 8 random downstream positions recover ≤ 0.035 . The conclusion sits on a few specific aggregator/delimiter tokens, not diffusely.

(iv) **Writability.** Overwriting an otherwise-consistent cache’s downstream notes with the *opposite* conclusion’s notes (the field token and prefix left intact and still implying the original answer) drives the decision to the injected conclusion: continuous recovery 0.99/0.98/1.02, with the top-8 note positions already flipping the belief in most instances. Editing (Section 4) is the benign use of this same write-access.

Off-template generalization. Re-running field-only vs. full-downstream recovery on (a) a 2-hop lookup, (b) free-form conversational phrasing, and (c) attribute lookup: field-only recovery is ≈ 0 for (a) and (b) across models (multi-hop -0.012 to $+0.006$; natural -0.012 to $+0.054$) with full-downstream 1.00, confirming the mechanism is not a template artifact; for near-verbatim attribute lookup (c) the field is partly a copy and carries 0.25–0.63, bounding the claim to *derived* conclusions.

Cross-family replication (Gemma-2, Mistral). To rule out a Qwen3/Llama tokenizer artifact we re-ran five probes (the locality probe and the four controls above) on **Gemma-2-9B** and **Mistral-7B** with a tokenizer-robust readout (space-prefixed single-token `cancel/deny`; Gemma-2’s soft-capping kept intact, as the attention-knockout hook used only by the original circuit-knockout probe otherwise corrupts it), on the gated `cancel/deny` task ($n=18$ primary, 18 dissociation each; Table 3). All five replicate: field-only ≈ 0 vs. full-downstream 1.0; dissociation trigger-only ≈ 0 vs. notes ≈ 1.0 ; top-8 $\gg$ random-8; injection ≈ 1.0 ; and write-before-read timing. Two honest notes: Mistral’s field-only recovery (0.137) is slightly above zero (still far below full-downstream 1.0), and random-8 specificity runs higher on this task (0.48–0.53 vs. ≈ 0.02 on the Qwen/Llama tool-call task) though top-8 (≈ 0.95) still dominates.

D A component-level circuit for memoized inference

Section 3 localizes the memoized conclusion (which tokens, which layers, a linear probe). This appendix pushes from *localization* to a *component-level* account—which attention heads, which di-

Table 3: Cross-family replication of the five deep-mechanism probes (Gemma-2-9B, Mistral-7B). Recovery toward the target conclusion; bootstrap means.

model	field-only	full-down	trigger-only	notes	top-8 / rand-8	write/commit depth
Gemma-2-9B	0.005	1.0	−0.00	1.0	0.96/0.48	0.26/0.48
Mistral-7B	0.137	1.0	0.001	0.995	0.94/0.53	0.19/0.47

Table 4: Deep mechanism controls (recovery toward the target conclusion; three models). Left: dissociation—patching the differing rule token vs. the downstream notes, field held identical. Right: specificity—top- k vs. random- k downstream positions, matched count.

(i) Dissociation			(iii) Specificity				
model	trigger only	notes	model	top-8	rand-8	top-16	rand-16
Qwen3-8B	−0.007	1.004	Qwen3-8B	0.78	0.009	0.92	0.06
Qwen3-4B	−0.007	1.009	Qwen3-4B	0.74	0.035	0.86	0.04
Llama-3.1-8B	+0.007	0.998	Llama-3.1-8B	0.79	0.005	0.86	0.04

rection, attention vs. MLP—with five interventions on the polarity 2×2 task, where the field is held *byte-identical* across the conclusion flip (only one rule-trigger token differs), so every “conclusion” signal we attribute cannot be field content. Primary model Llama-3.1-8B, replicated across *four families*—Qwen3-8B, Gemma-2-9B, Mistral-7B (Table 5)—all numbers are causal-intervention recoveries with bootstrap CIs over $n=12$ instances (Figure 16; scripts `esys/circ_*.py`, records `results/circ_*`).

(1) Named read and write heads (Figure 16a,f). We rank heads by direct attribution and confirm causally by patching a single head’s output (clean $\leftrightarrow$ corrupt) and re-reading the decision. *Read heads*—at the decision token—form a concentrated, nameable set: patching the top k jointly recovers 0.19/0.45/0.59/0.70/0.78 of the decision at $k=1/3/5/8/12$ on Llama-3.1-8B, and the top-12 reach 0.77/0.72/0.82 on Qwen3-8B/Gemma-2-9B/Mistral-7B (random-head control ≈ 0 on all four). The strongest read heads sit in late layers and attend decision $\rightarrow$ aggregator (Llama 26.3 at 0.46, Qwen 26.25 at 0.43, Gemma 26.9 at 0.47, Mistral 20.21 at 0.31). *Write heads*—at the aggregator, at prefill—are more distributed: single-aggregator patching of the top heads saturates at 0.28/0.08/0.24/0.11 (Llama/Qwen/Gemma/Mistral), because the write is spread over *many* aggregator tokens (the suffix-concentration of Section 3), so patching one position under-counts it. The decision-side read is thus a tight bottleneck; the write is diffuse.

(2) A causal conclusion direction (Figure 16b). A single difference-of-means direction $\hat{d}$ on the aggregator residual (fit *leave-one-scenario-out*, field-balanced) transfers the conclusion: injecting the $\hat{d}$ -component of the clean $\rightarrow$ corrupt residual into a corrupt run recovers 0.22–0.23 of the decision at Llama L12–14— $\approx 39\%$ of the full single-site residual patch and $\sim 25\times$ a random 1-D direction (0.01); the difference-of-means direction beats a logistic-probe direction (0.02), a known DAS phenomenon. So the conclusion has a real shared linear component, but it is *redundantly* coded (projecting $\hat{d}$ out of a clean run drops the decision only ~ 0.17). The direction is causal but similarly redundant on Gemma-2-9B (along 0.10 vs. random 0.00, of full 0.61) and Mistral-7B (along 0.07 vs. random 0.00, of full 0.36). On Qwen3-8B the conclusion is committed so strongly (logit gap ≈ 22 vs. ≈ 4) that single-site directional recovery is near zero in relative terms—the read-side circuit, which acts at the decision bottleneck, is the robust cross-family result.

Table 5: **The circuit replicates across four families.** Read = decision recovery from jointly patching the top-12 named read heads (control in parens). Write = top- k write-head recovery at the single top aggregator (distributed, so a floor). Scrub: faithfulness drift under same-conclusion resampling (want ≈ 0) and note-vs-rest interchange recovery. Attn = attention share of the write at the causal layer. Dir = 1-D conclusion-direction recovery (along $\hat{d}$ vs. random).

model	read (top-12)	write	scrub: drift / note / rest	attn share	dir: along / rand
Llama-3.1-8B	0.78 (0.00)	0.28	0.04 / 0.79 / 0.19	0.60–0.62	0.22 / 0.01
Qwen3-8B	0.77 (0.00)	0.08	0.04 / 0.68 / 0.32	0.44–0.56	— (over-committed)
Gemma-2-9B	0.72 (0.00)	0.24	0.02 / 0.82 / 0.18	0.82–0.92	0.10 / 0.00
Mistral-7B	0.82 (0.00)	0.11	0.03 / 0.71 / 0.30	0.42–0.66	0.07 / 0.00

(3) **A sparse SAE feature: decode $\neq$ cause (Figure 16e).** We train a TopK SAE (16k dict, $k=32$, FVU ≈ 0.001) on layer-14 residuals over diverse prompts. The conclusion is *sparsely decodable*: two features reach AUC = 1.00 at separating SAFE from UNSAFE on the aggregator (field-controlled), with ~ 5 features above 0.95. Yet it is *causally distributed*: clamping the single best feature to its SAFE level in a UNSAFE run recovers ≈ 0 , while clamping the top 3/10/30 jointly recovers 0.30/0.50/0.54 (random-feature control ≈ 0 ; necessity 0.30–0.33). A feature can read the conclusion out perfectly without being causally sufficient alone—the feature-level echo of the decodability-vs-causation dissociation in Section 3.

(4) **Causal scrubbing of write $\rightarrow$ note $\rightarrow$ read (Figure 16d).** We test the hypothesis “the decision is a function of the conclusion carried by the aggregator note” by resampling activations from inputs that agree with that labelling (position-aligned KV transplants from same-scenario donors). *Faithfulness*: resampling the entire downstream from a *same-conclusion* different-input donor preserves the decision (drift 0.02–0.04 of the gap; decision-logit cosine 0.999–1.000) on both models. *Necessity*: resampling only the *note* from an *opposite-conclusion* donor flips the decision 0.79/0.68/0.82/0.71 (Llama/Qwen/Gemma/Mistral), whereas resampling the rest of the downstream moves it only 0.19/0.32/0.18/0.30. The note governs; everything else is interchangeable.

(5) **Attention writes the note (Figure 16c).** Decomposing the SAFE–UNSAFE aggregator residual into per-layer attention- and MLP-block contributions (exact: the embedding is identical across the pair), attention contributes the majority of the write at the causal conclusion layer—0.62/0.60 at Llama L12/L14, and a striking 0.82–0.92 on Gemma-2-9B—with Qwen3-8B (0.44–0.56) and Mistral-7B (0.42–0.66) closer to parity. Attention is the primary or co-primary writer that routes the field/rule information onto the aggregator in every family, with the MLPs contributing the remainder.

Summary. The picture is a *distributed write, concentrated read*: at mid layers, attention routes the field-conditioned conclusion onto aggregator tokens (Exp. 5), in a redundant code that is sparsely decodable but causally spread over ~ 10 –30 features / a shared low-rank direction (Exp. 2–3); a small, nameable set of late read heads then funnels it into the decision logit (Exp. 1), and causal scrubbing confirms the note alone governs the decision (Exp. 4). This connects the “models take notes” account to the head-level circuits of Wang et al. [24] and to delimiter-token aggregation [11], while remaining a statement about the *KV cache* an inference system already stores.




Figure 16: **A component-level circuit for memoized inference** (deep dives Llama-3.1-8B; head, write, and scrubbing panels show all four families). (a) Cumulative decision recovery from jointly patching the top- k named *read* heads (concentrated, ~ 0.78) vs. *write* heads (distributed). (b) A leave-scenario-out difference-of-means *conclusion direction* causally transfers the decision far above a random 1-D direction. (c) At the causal conclusion layer, *attention* writes the majority of the note. (d) Causal scrubbing: same-conclusion resampling is faithful (drift ≈ 0); swapping the *note* to the opposite conclusion flips the decision while swapping the rest does not. (e) An SAE feature *decodes* the conclusion perfectly (AUC= 1.0) yet is not causally sufficient alone—the cause is spread over ~ 10 – 30 features. (f) Read heads with the largest causal effect are those that attend decision $\rightarrow$ aggregator.

E Editing: the baseline frontier and the K-sweep

Figure 17 plots the cost/correctness frontier behind Section 4: there is no single dominant method. *field+erratum* and the bare *ERRATUM* reach full correctness at ~ 5 – 13% recompute with no prompt surgery; *hoist-to-end* matches them but rewrites the prompt; the in-place edit and a KV-deviation-ranked *CacheBlend*-style recompute are cheap but incorrect on these gated decisions. Numeric values are in Table 6. Figure 18 gives the full *field+selective@K* sweep across seven models, making the model-dependence of the minimal recompute explicit: small K suffices for the Qwen3 family but not for several others—the tool is effective but unreliable.

Figure 19 shows that the editing fix is an *attention-architecture* method: the erratum recovers the decision under reasoning on attention (GQA) and sliding-window backbones, partially on a hybrid attention+SSM model, and weakly on a pure SSM whose recurrent state has no per-token look-back.




Figure 17: Cost/correctness frontier.




Figure 18: field+selective@K across models.

Table 6: Baseline frontier (Qwen3-8B, 8 gated tasks). “surgery” = requires rewriting the prompt.

method	P(correct)	recompute	prompt surgery
full repreload	1.00	100%	no
hoist-to-end	1.00	5.2%	yes
field+erratum	1.00	12.6%	no
ERRATUM	1.00	12.0%	no
CacheBlend@15%	0.13	14.5%	no
in_place edit	0.00	0.6%	no
stale (reuse)	0.00	0%	no

F Composing: per-domain scorecards, multimodal, and the MLA adapter

Figure 20 is the composable scorecard: decision agreement between a transplanted skill and full recompute, by model and content type (facts in two insertion points, and agentic tool-calling). Standard attention models are at or near 1.0; the sliding-window Gemma models are the consistent exception (addressed in Section 8). Figure 22 breaks the multimodal result down by task category, and Figure 21 reports the MLA decoupled- k_{pe} adapter fidelity and the transplant TTFT speedup across models. Figure 23 diagrams the three attention-variant adapters of Section 8: what each representation caches, and exactly which slice of it the reposition touches.

G Long-horizon robustness and the cross-referential memory test

No compounding error over a long trajectory (Figure 24). A standing risk for any leave-stale scheme is that small per-edit errors *compound* over a long agent trajectory. We test it directly. A single gated field (a clearance level) toggles between a granting and a denying value every turn over a 28-turn trajectory; we maintain ONE evolving KV cache—reuse the static prefix forever, apply each state change as an appended erratum, never recompute the downstream—and compare its per-turn decision to a *full repreload of the byte-identical token sequence* (errata included). The only difference between the two is whether downstream KV was recomputed after each change, so per-turn agreement isolates exactly the cost of leaving KV stale as a function of trajectory length. Across three families (Qwen3-8B, Llama-3.1-8B, Mistral-7B) the next-token *decision logits stay faithful*—cosine 0.987–0.999 with a flat first-third→last-third profile (e.g. 0.992 → 0.996 on Llama)—and the agreement-vs-turn slope is within $\pm 0.01/\text{turn}$ (no systematic decline). Discrete decision agreement is high in aggregate (0.79–0.99) but noisier than the cosine, because these gated decisions sit near the




Figure 19: Erratum recovery under reasoning across attention, sliding-window, hybrid, and pure-SSM backbones.




Figure 20: Composable agreement by model $\times$ content type.




Figure 21: MLA adapter fidelity (a) and transplant TTFT speedup across models (b).

action boundary (oracle accuracy 0.52–0.82), where a sub-percent logit difference can flip a discrete choice—the same boundary sensitivity noted for greedy CoT in Section 7. The leave-stale+erratum cache does not degrade with trajectory length: there is no compounding drift, only boundary noise.

Cross-referential facts (the Section 7 E4 test). E4 found splitting *independent* relevant facts across independently-precompiled blocks decision-lossless. We test the case it left open—a genuinely *cross-referential* chain. A memory contains a DEFINITION line A naming which setting governs the request, whose value B lives elsewhere; the decision needs the two-hop resolution $A \rightarrow B$. We lay A and B so they straddle a block boundary and compare a transplant that *splits* them (boundary between A and B) against one that keeps them *colocated* (boundary moved past both), holding everything else fixed. On Llama-3.1-8B—which performs this two-hop decision at full-recompute accuracy 0.85 ($n=80$, balanced)—splitting drops agreement with full recompute to 0.46 versus 0.76 colocated, a 0.30 penalty (25 vs. 1 discordant instances, McNemar $p=8 \times 10^{-7}$); an *independent* fact pair split the same way costs only 0.04 (Mistral-7B). The mechanism explains it: block B precompiled in isolation never attended to its referent A , so the chain is not memoized and the splice cannot restore it. (Qwen3-8B is at chance on the direct two-hop decision—consistent with Section 7’s observation that direct gated decisions are at chance for the reasoning-native family $\leq 32B$ —so it is uninformative here.) The guidance is simple: keep cross-referential facts within one




Figure 22: Image-KV transplant by task category (perception / reasoning / agentic), full re-encode vs. transplant, across four vision-language models. The transplant tracks full accuracy category-by-category.

precompiled block; independent facts may be split freely.

H The user-memory agent


The agent of Section 7 keeps the layout [system] [trajectory] [MEMORY] [query]. The system prompt is prefilled once; the trajectory grows by cached deltas; the user-memory chunk is precompiled once in isolation and RoPE-repositioned to float just before the query each turn (an $O(L_{\text{mem}})$ re-rotation, no re-prefill) with one boundary token repaired; a memory change is applied either by recompiling the isolated chunk ($O(L_{\text{mem}})$, once) or by appending a salient erratum into the cached trajectory stream (append-only, composes with prefix caching). Confirmatory protocol: gated decisions whose governing facts live in a Markdown memory, balanced labels, chain-of-thought as the competent regime (direct one-shot memory-gated decisions are at chance for $\leq 8\text{B}$ models), with cluster-bootstrap CIs (10^4 , persona-level), TOST equivalence ($\delta=0.03$; cosine ≥ 0.98), GEE-logistic (cluster-robust), McNemar, and BH-FDR. Faithfulness is read against a token-matched full-reprefill oracle. Pre-registered hypotheses, margins, and the inclusion gate (oracle accuracy ≥ 0.80) are released with the code.

I Online serving and weight-editing methodology


Online vLLM serving (Figure 15a). vLLM V1 AsyncLLMEngine, CUDA graphs enabled (*not* enforce_eager), continuous batching, automatic prefix caching on, bfloat16, GPU memory utilization 0.85. Workload: a shared $\sim 8,066$ -token policy (real τ^2 -bench retail policy plus neutral padding) with one mutable field; 96 requests per arm, 64 output tokens each. Requests arrive as a *Poisson* process at offered rates $\{2, 4, 8, 16\}$ req/s and unthrottled (saturation). We timestamp first-token and completion per request in the client loop (TTFT, TPOT, end-to-end) and read each arm’s APC hit-rate from the engine’s Prometheus counters (vllm:gpu_prefix_cache_hits, ..._queries). The in-prefix baseline writes the new field value early (invalidating downstream APC blocks); the erratum keeps the old prefix and appends the update. Headline: throughput speedup grows $1.6 \times \rightarrow 14.5 \times$ as load rises to saturation; p_{90} TTFT 53–398 $\times$ lower; APC hit-rate 98.5% vs. 1.0%.

Weight editing (Table 1). ROME is implemented from scratch (uncentered key covariance $C = \mathbb{E}[kk^\top]$ at a mid MLP layer over a text sample, an optimized value v^* , and the closed-form rank-one update of down_proj) and *validated on the canonical factual edit* (“the Eiffel Tower is

(a) **MLA adapter: re-rotate only the decoupled RoPE sub-vector k^{pe} ; the latent c_t is position-free**



(b) **M-RoPE adapter (vision): re-rotate only the temporal axis t ; spatial h, w are intrinsic to the image**



(c) **Sliding-window fix (Gemma): keep full per-token KV; let the attention mask enforce the window**




Figure 23: The attention-variant adapters. (a) MLA caches a position-free latent c_t plus a small decoupled-RoPE sub-vector k_{pe} ; repositioning re-rotates only k_{pe} and copies the latent as-is. (b) M-RoPE factors the rotary channels into temporal/height/width axes—sectioned (Qwen2.5-VL) or interleaved (Qwen3-VL); moving an image within a trajectory re-rotates only the temporal channels, since the spatial axes are intrinsic to the image. (c) Sliding-window layers default to a window-truncated cache, which breaks any splice past W ; keeping the full per-token KV and letting the attention mask enforce the window restores uniform edit/splice semantics.

in” *Paris*→*Rome*, locality preserved) before use, so the baseline is faithful. LoRA fine-tunes ($r=8$, q, v , down projections) on the stale-context decision until it flips. All methods are evaluated on the same gated decision; cross-request contamination is measured over 8 held-out orders that are genuinely still pending (correct = cancel), and collateral over a 10-item battery of unrelated single-token gated decisions.

J Reproducibility

All numbers are from local runs; result records and figure-generation scripts are released with the code at <https://github.com/bojieli/programmable-kv>; an interactive companion is at <https://01.me/research/programmable-kv/>. Bootstrap confidence intervals use 10^4 resamples. The unified-agent study uses 10 domains $\times$ 10 instances (300 decisions) per model, except the two Gemma models (40 instances, 120 decisions); the multimodal and agentic studies use $N=120$ and $N=108$. The cross-family replication, online serving, and weight-editing studies ship with their result JSONs (scripts under `esys/`: `mechd_replicate.py`, `vllm_serving_online.py`, `rome.py`, `weight_editing_compare.py`).




Figure 24: **No compounding error over a long trajectory.** A gated field toggles every turn for 28 turns; one evolving leave-stale+erratum cache vs. full refill of the identical text. (a) Per-turn decision agreement stays high with boundary noise (no downward trend). (b) The decision-logit cosine stays flat at 0.99+—no drift with trajectory length.